

Business Programming for E&BI

Christoph Walsh

What is R and RStudio?

- R is a *programming language* which specializes in statistical computing and graphics.
 - ▶ A programming language is a way to instruct a computer to perform operations via written text.
 - ▶ When programming, we need to be very exact. Otherwise the computer will throw an error or do something we didn't intend it to do!
- RStudio is a desktop application where you can write R code, execute R programs, and view plots created by R.

Why Learn R?

R has many benefits over some alternatives:

- R is free and open source.
- Large active community creating packages and providing support.
- Easier to learn than some alternatives.
- Availability of RStudio for free.
- R is extremely versatile. For example, these slides were made in RStudio!

Your R journey also does not end here: you will also use R in later courses (e.g. Quantitative Methods courses).

Installing R

- To download and install R, go to <https://mirror.lyrahosting.com/CRAN>.
- To download and install RStudio, go to <https://posit.co/download/rstudio-desktop/>.

The R Console

- We can perform calculations in the Console tab in RStudio.
- At the most basic level, we can use it as a simple calculator:
 - ▶ Add: $2 + 3$
 - ▶ Subtract: $5 - 3$
 - ▶ Multiply: $2 * 3$
 - ▶ Divide: $3 / 2$
 - ▶ Exponentiation: 2^3
 - ▶ Combining operations: $(2 + 4) / (4 * 2)$


R Functions

- Just like Excel, R has functions.
- These functions work in a very similar way:
 - ▶ Functions have names, and we provide *arguments* to functions inside parentheses.
- In the next few slides we will see some examples of mathematical functions and how to evaluate them in R.

Absolute Value

- The absolute value function turns negative numbers into positive ones and has no effect on zero or positive numbers.

$$|x| = \begin{cases} x & \text{if } x \geq 0 \\ -x & \text{otherwise} \end{cases}$$



Absolute Value in R

- The absolute value function in R is called `abs`. We can use it as follows:

```
abs(-2)
```

```
[1] 2
```


```
abs(3)
```

```
[1] 3
```

- We can see help about a function using `help(abs)` or `?abs`

Principal Square Roots

- The square root of a number is the y that solves $y^2 = x$.
- If $x = 4$, both $y = -2$ and $y = 2$ solve $y^2 = x$.
- The principal square root is the positive y solving this.



Principal Square Roots in R

- We can use the `sqrt()` function or take to the power of $\frac{1}{2}$ to take the square root.

```
sqrt(9)
```

```
[1] 3
```

```
9^(1/2)
```

```
[1] 3
```

- The cubed root is the y that solves $y = x^3$. We can calculate this in R by taking the power of $\frac{1}{3}$.

```
8^(1/3)
```


```
[1] 2
```

The Exponential Function

The exponential function is a very common function in statistics:

$$e^x = \lim_{n \rightarrow \infty} \left(1 + \frac{x}{n}\right)^n$$

Note: you don't need to know this function for the exam.



The Exponential Function in R

- In R we use the `exp()` function to calculate the exponential of a number:

```
exp(0)
```

```
[1] 1
```

```
exp(1)
```


```
[1] 2.718282
```

The Logarithm

- Another common mathematical function is the logarithm, which is like the reverse of exponentiation.
- The log of a number x to a base b , denoted $\log_b(x)$, is the number of times we need to multiply b by itself to get x .
- For example, $\log_{10}(100) = 2$, because $10 \times 10 = 100$. We need to multiply the base $b = 10$ by itself twice to get to $x = 100$.

The Natural Logarithm

- A special logarithm is the natural logarithm, $\log_e(x)$, which is the logarithm to the base $\exp(1) = e^1 \approx 2.7183$. This is also written as $\ln(x)$.



The Logarithm in R

- In R we use the `log()` function to calculate the natural logarithm:

```
log(1)
```

```
[1] 0
```

- We can calculate the logarithm of a number to a different base using the `base` argument. For example, for $\log_{10}(100)$:

```
log(100, base = 10)
```

```
[1] 2
```

The Assignment Operator: <-

- We can store objects using the *assignment operator*, <-.
- For example:

```
a <- 2  
b <- 3
```

- a and b are then visible in the **Environment** tab in RStudio.
- We can then use a and b for calculations:

```
a + b  
  
[1] 5
```

Common Object Types: Numerical, Logical and Character

- Numerical vectors (list of numbers):

```
a <- c(1, 3, 7, 2)
```

- Logical vectors (list of Yes/No responses):

```
a <- c(TRUE, FALSE, TRUE, TRUE)
```

- Character vectors (list of letters/words):

```
a <- c("abc", "def", "ghi")
```

Common Object Types: Factors

- Categorical variables are stored in R as “factors”.
- For example, imagine a survey asking how long it took for people to get to campus (in minutes) and their travel mode (one of “cycle”, “train”, or “walk”).
- You could store these variables as a numerical and a character variable:

```
time <- c(25, 20, 15, 10, 17, 30)
travel_mode <- c("train", "train", "walk",
                 "cycle", "walk", "train")
```

- But if we tell R that this travel mode variable is a categorical variable, it will be useful for operations that we will learn later. We can convert this to a factor using the factor function:

```
travel_mode <- factor(travel_mode)
```

Common Object Types: Data Frames and Lists

- A `data.frame` collects vectors of the same length into a single dataset.
- We can convert the `time` and `travel_mode` variables into a `data.frame` as follows:

```
df <- data.frame(travel_mode, time)
```

- We can also “View” it in RStudio by clicking on it in the Environment tab.
- A `data.frame` is actually a special type of list:

```
my_list <- list(x = 1:3, y = TRUE, z = c("a", "b"))
```

- While vectors must have all elements of the same type (numeric/logical/character/factor), lists can have elements of any type *and any length*.
- Dataframes can have columns of different types, but all columns must have the same length.

Indexing Vectors with Numbers

- The elements of a vector like:

```
a <- c(1, 2, 4, 3, 2)
```

are indexed 1-5.

- We can extract elements from this vector using these indices:

```
a[3]
```

```
[1] 4
```

```
a[c(1, 3, 4)]
```

```
[1] 1 4 3
```

Indexing Vectors with Logical Vectors

- We can also index vectors using an equal-length logical vector, where we extract only the elements that are TRUE:

```
a <- c(1, 2, 4, 3, 2)
a[c(TRUE, FALSE, TRUE, TRUE, FALSE)]
```

```
[1] 1 4 3
```

Sequences

We can create vectors that are sequences of numbers in different ways:

```
1:10
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

```
10:1
```

```
[1] 10 9 8 7 6 5 4 3 2 1
```

```
seq(from = 10, to = 100, by = 10)
```

```
[1] 10 20 30 40 50 60 70 80 90 100
```

```
seq(from = 0, to = 1, length.out = 5)
```

```
[1] 0.00 0.25 0.50 0.75 1.00
```

Repeating Numbers and Vectors

To save time typing, we can repeat numbers and vectors with the `rep()` function:

```
rep(1, times = 20)
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

```
rep(1:3, times = 4)
```

```
[1] 1 2 3 1 2 3 1 2 3 1 2 3
```

```
rep(1:3, each = 4)
```

```
[1] 1 1 1 1 2 2 2 2 3 3 3 3
```

Summary Statistics for Vectors

```
a <- 1:10
```

```
a
```

```
[1]  1  2  3  4  5  6  7  8  9 10
```

```
length(a)
```

```
[1] 10
```

```
min(a)
```

```
[1] 1
```

```
max(a)
```

```
[1] 10
```

Summary Statistics for Vectors

```
mean(a)
```

```
[1] 5.5
```

```
median(a)
```

```
[1] 5.5
```

```
sum(a)
```

```
[1] 55
```

```
summary(a)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1.00	3.25	5.50	5.50	7.75	10.00

Tabulate a Vector

```
a <- c(1, 3, 2, 4, 4, 2, 4)
table(a)
```

a

1 2 3 4

1 2 1 3

Example: A Linear Mortgage

- You take out a linear mortgage for €150,000 with a 15-year term and an interest rate of 7%. How much interest will you pay on this loan?
- You pay $€150,000/15=€10,000$ per year off the loan principal.
- The debt follows a sequence 150,000, 140,000, 130,000, ..., 10,000, 0.
- The interest of 7% each period is on the remaining debt.
- We can get the total interest payments with R using:

```
debt <- seq(from = 150000, to = 0, by = -10000)
interest <- debt * 0.07
sum(interest)
```

```
[1] 84000
```

Example: Sum of Geometric Sequences

- The parents of a child choose to invest €1,250 at 5% interest each year from when the child is born until (and including) their 20th birthday. On the 20th birthday, how much will the investment be worth?
- The first investment is worth: $1250 \times (1.05)^{20} = 3316.62$.
- The second investment is worth: $1250 \times (1.05)^{19} = 3158.69$.
- The last investment won't have received interest and is worth $1250 \times (1.05)^0 = 1,250$.

```
sum(1250 * 1.05^(20:0))
```

```
[1] 44649.06
```

This is the same as using the solution to the sum of a geometric sequence:

```
1250 * (1.05^21 - 1) / 0.05
```

```
[1] 44649.06
```

Comparing Numerical Vectors

```
a <- 1:5
```

```
b <- 5:1
```

```
a
```

```
[1] 1 2 3 4 5
```

```
b
```

```
[1] 5 4 3 2 1
```

```
a < b
```

```
[1] TRUE TRUE FALSE FALSE FALSE
```

```
a <= b
```

```
[1] TRUE TRUE TRUE FALSE FALSE
```

Comparing Numerical Vectors

```
a == b
```

```
[1] FALSE FALSE TRUE FALSE FALSE
```

```
a != b
```

```
[1] TRUE TRUE FALSE TRUE TRUE
```

```
a >= b
```

```
[1] FALSE FALSE TRUE TRUE TRUE
```

```
a > b
```

```
[1] FALSE FALSE FALSE TRUE TRUE
```

Comparing Logical Vectors

```
a <- c(TRUE, TRUE, FALSE, FALSE)
b <- c(TRUE, FALSE, TRUE, FALSE)
```

We use `&` for *logical AND*, `|` for *logical OR* and `!` for *logical NOT*.

```
a & b
```

```
[1] TRUE FALSE FALSE FALSE
```

```
a | b
```

```
[1] TRUE TRUE TRUE FALSE
```

```
!a
```

```
[1] FALSE FALSE TRUE TRUE
```

R Scripts

- When working on a project, it's often better to write all the commands you want to run in an *R script* instead of directly into the console:
 - ▶ It documents and saves your work.
 - ▶ It makes your work shareable and reproduceable.
 - ▶ You can edit earlier commands in a chain of commands.
 - ▶ It's easier to spot mistakes.
- There are different ways of “running” an R script:
 - ▶ Selecting all or a subset of lines and hitting “Run”.
 - ▶ Sourcing “with echo” and “without echo”.

Commenting in R Scripts

- When writing code it's good practice to write “comments” to help others (and you!) understand your code.
- Anything written after a # symbol is ignored by the R console. So we precede all comments with a #.

```
# Set values of a and b:  
a <- 2  
b <- 3  
print(a + b) # Compute the sum of a and b and print:
```

```
[1] 5
```

- Commenting is also useful if you want to temporarily disable a certain part of your script. You just need to put a # before the commands you want to disable.
 - ▶ This is called “commenting out”.
 - ▶ Select the lines you want to comment out, then go to Code → Comment/Uncomment Lines.

CSV files

- The most common way we read data into R is through CSV (comma-separated values) files.
- These are plain text files with a .csv extension.
- We would store our travel mode survey data in a CSV file like this:

```
travel_mode,time  
train,25  
train,20  
walk,15  
cycle,10  
walk,17  
train,30
```

CSV files: Commas in data points

- Each line needs to have the same number of commas.
- If data points contain commas themselves (e.g. suppose a category was `train, cycle`), we need to wrap the data points with quotes.
- So often you will see CSV files where the text is wrapped in quotes:

```
"travel_mode","time"  
"train",25  
"train",20  
"walk",15  
"cycle",10  
"walk",17  
"train",30
```

CSV files: Comma decimal separators

- In continental Europe, commas are used as decimal separators: “one and a half” is written as 1,5.
- Clearly this will cause problems with CSV files!
- So sometimes you might see files with ; delimiting variables instead of ,.
- An example of that would be:

```
"travel_mode";"time"  
"train";25,0  
"train";20,0  
"walk";15,0  
"cycle";10,0  
"walk";17,0  
"train";30,0
```

- When this happens, we need to tell R that the file is using a ; separator.
- In the exam, however, we'll deal exclusively with more standard CSV files.

Reading CSV files into R

- In order to read a CSV file into R, we need to tell R where the file is located on our computer.
- There are three different approaches to do this:
 - 1 The absolute path method
 - 2 The relative path method
 - 3 The RStudio Projects method (my recommendation!)
- For this we save the travel mode data in a file called `test.csv`.

Reading CSV files into R: The Absolute Path Method

- This approach involves giving R the full path to the file.
- On Windows, the full path of would look something like `C:\Users\username\Documents\test.csv`.
- However, the backslash has a special purpose in R so we can't use this.
- We need to use either forward slashes ("/") or *double-backslashes* ("\\"):
 - ▶ `C:/Users/username/Documents/test.csv`
 - ▶ `C:\\Users\\username\\Documents\\test.csv`
- The fastest way to get the full path to a file is with the `file.choose()` command.
- Run the command, navigate to the file, and then the full file path with appear in the console. You can then copy this to your clipboard.
- You can then paste this as an argument into the `read.csv()` command:

```
df <- read.csv("C:\\Users\\username\\Documents\\test.csv")
```

- When we use this approach, it doesn't matter what the current working directory of the R process is.

Reading CSV files into R: The Relative Path Method

- We can find out what the current working directory is with `getwd()`.
- To change it to the Documents folder, we can use the `setwd()` command.
- When the R process is in the working directory with the data, we only need to provide the name of the file in the `read.csv()` command.
- The steps would then be:

```
setwd("C:\\Users\\username\\Documents\\")  
df <- read.csv("test.csv")
```

- Suppose the full path of the file was actually
"C:\\Users\\username\\Documents\\data\\test.csv"
- We could then read in the data from the Documents folder by only giving the *relative path* to the file, which is "data\\test.csv":

```
setwd("C:\\Users\\username\\Documents\\")  
df <- read.csv("data\\test.csv")
```

Reading CSV files into R: The RStudio Projects Method

- If you share your code with someone, they will have to edit these lines that read in the data, or change the working directory. This is not ideal!
- A better way to deal with file paths is by using the RStudio Project feature.
- Suppose you saved your data in a folder called PQS on your computer.
- To go File → New Project, choose “Existing Directory”, and navigate to the PQS folder.
- When you are in the PQS project, RStudio automatically changes the current working directory to that folder. Then you don't need to provide the full file path, or use the `setwd()` command.
- Therefore I recommend this approach over the previous methods.

R Packages: Installing a Package

- Up to now, all the functions we have been using come by default with R.
 - ▶ We call the default functionality “base R”.
- But people have written *packages* that expand the functionality of R to do more things.
- For example, base R is not able to read in data from an Excel file.
 - ▶ But there are several packages that can do this.
 - ▶ We'll learn how to do this using the `readxl` package.
- To install a package, you can use the command:

```
install.packages("readxl")
```

- Alternatively, you can install the package using Tools → Install Packages... in RStudio.

R Packages: Loading and Using a Package

- R doesn't load up the functions of all the installed packages by default. We also need to *load* a package after we install it.
- We can do this with the `library()` command:

```
library(readxl)
```

- If a package is not installed, the `library` function will return an error.
- To read in a file called `test.xlsx` in the current working directory, we do:

```
df <- read_excel("test.xlsx")
```

- The `read_excel()` function loads the data as a `tibble`, which is like a `data.frame` but with a few extra features. We can force the data to be a plain `data.frame` with:

```
df <- data.frame(read_excel("test.xlsx"))
```

Dataframes: Eredivisie Data

	team	wins	draws	losses	goals_for	goals_against
1	AZ	20	7	7	68	35
2	Ajax	20	9	5	86	38
3	Excelsior	9	5	20	32	71
4	FC Emmen	6	10	18	33	65
5	FC Groningen	4	6	24	31	75
6	FC Twente	18	10	6	66	27
7	FC Utrecht	15	9	10	55	50
8	FC Volendam	10	6	18	42	71
9	Feyenoord	25	7	2	81	30
10	Fortuna Sittard	10	6	18	39	62
11	Go Ahead Eagles	10	10	14	46	56
12	NEC	8	15	11	42	45
13	PSV	23	6	5	89	40
14	RKC Waalwijk	11	8	15	50	64
15	SC Cambuur	5	4	25	26	69
16	Sparta Rotterdam	17	8	9	60	37
17	Vitesse	10	10	14	45	50
18	sc Heerenveen	12	10	12	44	50

Dataframes: Indexing

- We can get the 2nd row and 3rd column of df with `df[2, 3]`.
- We can the team name and number of wins variables for Ajax, Feyenoord and PSV with `df[c(2, 9, 13), c(1, 2)]`.
- We can extract all variables for Ajax with `df[2, ]`.
- We can extract all values for “goals for” with `df[, 5]`.
- We can also extract all values for a column with a variable name:
 - ▶ `df$goals_for`.
 - ▶ `df[, "goals_for"]`.
 - ▶ `df[["goals_for"]]`.
- We can extract several columns with:
 - ▶ `df[, c("team", "goals_for")]`
- To get the rows for all teams with at least 20 wins: `df[df$wins >= 20, ]`.

Dataframes: Creating Variables - Goal Difference

```
df$goal_diff <- df$goals_for - df$goals_against  
head(df)
```

	team	wins	draws	losses	goals_for	goals_against	goal_diff
1	AZ	20	7	7	68	35	33
2	Ajax	20	9	5	86	38	48
3	Excelsior	9	5	20	32	71	-39
4	FC Emmen	6	10	18	33	65	-32
5	FC Groningen	4	6	24	31	75	-44
6	FC Twente	18	10	6	66	27	39

Dataframes: Creating Variables - Total Points

$$\text{Points} = 3 \times \text{Wins} + 1 \times \text{Draws} + 0 \times \text{Losses}$$

```
df$total_points <- 3 * df$wins + df$draws  
head(df[, c("team", "wins", "draws", "losses", "total_points")])
```

	team	wins	draws	losses	total_points
1	AZ	20	7	7	67
2	Ajax	20	9	5	69
3	Excelsior	9	5	20	32
4	FC Emmen	6	10	18	28
5	FC Groningen	4	6	24	18
6	FC Twente	18	10	6	64

Dataframes: Creating Variables - Team Ranking

```
df <- df[order(df$total_points, df$goal_diff, decreasing = TRUE), ]  
df$ranking <- 1:nrow(df)  
head(df[, c("team", "total_points", "goal_diff", "ranking")])
```

	team	total_points	goal_diff	ranking
9	Feyenoord	82	51	1
13	PSV	75	49	2
2	Ajax	69	48	3
1	AZ	67	33	4
6	FC Twente	64	39	5
16	Sparta Rotterdam	59	23	6

Dataframes: Creating Variables - Relegation Status

```
df$relegation_status <- ""
df$relegation_status[df$ranking < 16] <- "No relegation"
df$relegation_status[df$ranking == 16] <- "Relegation playoffs"
df$relegation_status[df$ranking %in% 17:18] <- "Automatic relegation"

df[, c("team", "ranking", "relegation_status")]
```

	team	ranking	relegation_status
9	Feyenoord	1	No relegation
13	PSV	2	No relegation
2	Ajax	3	No relegation
1	AZ	4	No relegation
6	FC Twente	5	No relegation
16	Sparta Rotterdam	6	No relegation
7	FC Utrecht	7	No relegation
18	sc Heerenveen	8	No relegation
14	RKC Waalwijk	9	No relegation
17	Vitesse	10	No relegation
11	Go Ahead Eagles	11	No relegation
12	NEC	12	No relegation
10	Fortuna Sittard	13	No relegation
8	FC Volendam	14	No relegation
3	Excelsior	15	No relegation
4	FC Emmen	16	Relegation playoffs
15	SC Cambuur	17	Automatic relegation
5	FC Groningen	18	Automatic relegation

The %in% Operator

- When we write `a %in% b` we are checking for each element in `a` if there is a *matching* element somewhere in `b`.

```
a <- 1:6  
b <- c(3, 5, 7)  
a %in% b
```

```
[1] FALSE FALSE TRUE FALSE TRUE FALSE
```

An equivalent but longer way of doing the same thing would be:

```
a == b[1] | a == b[2] | a == b[3]
```

```
[1] FALSE FALSE TRUE FALSE TRUE FALSE
```

Dataframes: Summarizing a Dataframe

The `summary()` command gives the summary statistics of all variables in `df`:

```
summary(df[, c("team", "wins", "draws", "losses")])
```

team	wins	draws	losses
Length :18	Min. : 4.00	Min. : 4.000	Min. : 2.00
N.unique :18	1st Qu.: 9.25	1st Qu.: 6.000	1st Qu.: 7.50
N.blank : 0	Median :10.50	Median : 8.000	Median :13.00
Min.nchar: 2	Mean :12.94	Mean : 8.111	Mean :12.94
Max.nchar:16	3rd Qu.:17.75	3rd Qu.:10.000	3rd Qu.:18.00
	Max. :25.00	Max. :15.000	Max. :25.00

Dataframes: Peeking at your data

- We use `head(df, n = 4)` to see the first 4 rows of `df`. `head(df)` on its own shows 6 rows by default.
- We use `tail(df, n = 2)` to see the last 2 rows of `df`.

```
nrow(df) # see the number of rows in df
```

```
[1] 18
```

```
ncol(df) # see the number of columns in df
```

```
[1] 10
```

```
dim(df) # see both the number of rows and number of columns in df
```

```
[1] 18 10
```

Dataframes: See all variable names

- names(df) shows the variable names of all variables in df.

```
names(df)
```

```
[1] "team"           "wins"           "draws"
[4] "losses"         "goals_for"      "goals_against"
[7] "goal_diff"      "total_points"   "ranking"
[10] "relegation_status"
```

- Later we will learn how to use this command to change the names of variables in df.

Data Cleaning

Very often when we have a spreadsheet, we need to do some “cleaning” before we can work with it in R.

This can happen when:

- The data don't start at the top of the file because the first few rows contain some other information.
- The dates are not formatted correctly.
- Numbers are interpreted as characters.
- The data contain extra columns that we don't want.
- There are rows with missing data that we want to omit.
- The variable names are not what we want them to be.

ASML Stock Price Data

We will learn how to clean data with the `asml-trades.csv` data.

The variable names are:

- `Date`: The date the data from that row are from.
- `Open`: The opening price of the stock on that day.
- `High`: The highest price the stock traded at on that day.
- `Low`: The lowest price the stock traded at on that day.
- `Last`: The price of the last-traded stock at on that day.
- `Close`: The closing price of the stock on that day.
- `Number.of.Shares`: The number of shares traded that day.
- `Number.of.Trades`: The number of trades made that day.

Skipping Rows

If the data don't begin at the top of a file, you can tell R to ignore the empty rows with the skip option:

```
df <- read.csv("asml-trades.csv", skip = 3)
```

```
summary(df)
```

Date		Open	High	Low
Length	:521	Min. :394.7	Min. :408.2	Min. :375.8
N.unique	:521	1st Qu.:535.5	1st Qu.:545.5	1st Qu.:525.8
N.blank	: 0	Median :592.1	Median :597.4	Median :582.5
Min.nchar	: 8	Mean :589.2	Mean :597.5	Mean :579.6
Max.nchar	:10	3rd Qu.:645.5	3rd Qu.:652.5	3rd Qu.:636.1
		Max. :770.5	Max. :777.5	Max. :764.2
		NAs :6	NAs :6	NAs :6
Last		Close	Number.of.Shares	Number.of.Trades
Min.	:397.4	Min. :397.4	Length :521	Length :521
1st Qu.	:535.9	1st Qu.:535.9	N.unique :516	N.unique :512
Median	:589.4	Median :589.4	N.blank : 0	N.blank : 0
Mean	:588.4	Mean :588.4	Min.nchar: 4	Min.nchar: 4
3rd Qu.	:644.0	3rd Qu.:644.0	Max.nchar: 7	Max.nchar: 6
Max.	:770.5	Max. :770.5		
NAs	:6	NAs :6		
Print.table				
Mode:logical				
NAs :521				

Formatting Dates

The date variable was read in as a character instead of a date.

```
head(df$Date, n = 3)
```

```
[1] "31/8/2021" "1/9/2021"  "2/9/2021"
```

The dates are in the format “dd/mm/yyyy”. We can convert these to dates using the `as.Date()` function, telling R the format the dates are in:

```
df$Date <- as.Date(df$Date, format = "%d/%m/%Y")
```

```
summary(df$Date)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
"2021-08-31"	"2022-03-01"	"2022-08-30"	"2022-08-29"	"2023-02-28"	"2023-08-29"

Formatting Dates: More Examples

```
as.Date("12/31/2023", format = "%m/%d/%Y")
```

```
[1] "2023-12-31"
```

```
as.Date("31-12-2023", format = "%d-%m-%Y")
```

```
[1] "2023-12-31"
```

```
as.Date("31/12/23", format = "%d/%m/%y")
```

```
[1] "2023-12-31"
```

```
as.Date("31 Dec 2023", format = "%d %b %Y")
```

```
[1] "2023-12-31"
```

```
as.Date("31 December 2023", format = "%d %B %Y")
```

```
[1] "2023-12-31"
```

Converting Characters to Numbers

- Although most values are numbers, there are some elements with "None" in the variables `Number.of.Shares` and `Number.of.Trades`.
- These character elements force the entire variable to be character, because all elements in a vector must have the same type.
- If we convert the character elements to NA (missing values), we can convert the variable to a numeric vector:

```
df$Number.of.Shares[df$Number.of.Shares == "None"] <- NA  
df$Number.of.Shares <- as.numeric(df$Number.of.Shares)
```

- We could skip the step of converting character elements to NA, but then R would warn us: NAs introduced by coercion.

Deleting Variables

We can delete variables by assigning NULL to that variable.

```
df$Print.table <- NULL
```

Alternatively we can drop variables using the column index of the variables we want to drop. We can drop the 9th variable (Print.table) with:

```
df <- df[, -9]
```

Dropping Rows with Missing Data

- For all variables apart from the date we having missing data on 6 rows. Let's take a look at which dates these are:

```
df$Date[is.na(df$Open)]
```

```
[1] "2022-04-15" "2022-04-18" "2022-12-26" "2023-04-07" "2023-04-10"  
[6] "2023-05-01"
```

- These are dates around Easter and Christmas when the stock market is closed.
- We can drop rows with any missing variables with the `na.omit()` function.

```
nrow(df)
```

```
[1] 521
```

```
df <- na.omit(df)  
nrow(df)
```

```
[1] 515
```

Renaming Variables

We can change a variable name using its column index as follows:

```
names(df)[7] <- "num_shares"
```

We can change a variable name using its current name as follows:

```
names(df)[names(df) == "Number.of.Trades"] <- "num_trades"
```

We can change multiple variable names at once with:

```
names(df)[2:5] <- c("open", "high", "low", "last")
```

Renaming Variables: Converting to Lower Case

We can convert characters to lower case with the `tolower()` function:

```
test <- c("hello!", "HELLO!", "Hello!", "HeLlO!")  
tolower(test)
```

```
[1] "hello!" "hello!" "hello!" "hello!"
```

We can convert all variable names to lower case with:

```
names(df) <- tolower(names(df))  
names(df)
```


```
[1] "date"          "open"          "high"          "low"           "last"  
[6] "close"         "num_shares"    "num_trades"
```

Introduction to Plotting

- We will now learn how to plot data with R.
- We will learn how to make:
 - ▶ Histograms: these display the distribution of a numeric variable.
 - ▶ Bar charts: these display frequencies of values for numeric or categorical data.
 - ▶ Scatter plots: these display the relationship between two numeric variables.
- There are two ways to make these plots:
 - 1 The “base R” approach (using built-in plotting functions in R)
These are quick and easy to make, but don’t look very nice.
 - 2 The “ggplot” approach (using the ggplot2 package).
These require more code, but are prettier and more customizable.

The Palmer Penguins Dataset

- We will use this famous dataset to demonstrate plotting.
- It contains the weight, gender, flipper length, and bill length and depth for 3 types of penguins: the adelie, the chinstrap and gentoo.



We can load the dataset with:

```
install.packages("palmerpenguins")  
library(palmerpenguins)  
data(penguins)
```

Summarizing the Data

```
summary(penguins)
```


species	island	bill_length_mm	bill_depth_mm
Adelie :152	Biscoe :168	Min. :32.10	Min. :13.10
Chinstrap: 68	Dream :124	1st Qu.:39.23	1st Qu.:15.60
Gentoo :124	Torgersen: 52	Median :44.45	Median :17.30
		Mean :43.92	Mean :17.15
		3rd Qu.:48.50	3rd Qu.:18.70
		Max. :59.60	Max. :21.50
		NAs :2	NAs :2

flipper_length_mm	body_mass_g	sex	year
Min. :172.0	Min. :2700	female:165	Min. :2007
1st Qu.:190.0	1st Qu.:3550	male :168	1st Qu.:2007
Median :197.0	Median :4050	NAs : 11	Median :2008
Mean :200.9	Mean :4202		Mean :2008
3rd Qu.:213.0	3rd Qu.:4750		3rd Qu.:2009
Max. :231.0	Max. :6300		Max. :2009
NAs :2	NAs :2		

Histograms in Base R

```
hist(penguins$body_mass_g)
```

Histogram of penguins\$body_mass_g



Bar Plots in Base R

The `table()` function applied to a vector shows the number of times each value appears.


```
table(penguins$species)
```

Adelie	Chinstrap	Gentoo
152	68	124

We can use a bar plot to visualize these relative frequencies.


Bar Plots in Base R

```
barplot(table(penguins$species))
```



Scatter Plots in Base R

```
plot(penguins$bill_length_mm, penguins$flipper_length_mm)
```




ggplot: The Grammar of Graphics

- Although it's possible to customize the plots from base R, we will instead learn how to produce nicer plots using the `ggplot2` package.
- The “gg” in `ggplot` stands for “Grammar of Graphics”, which is a scheme to *layer* elements in a plot.
- With the `ggplot2` package, we create plots by adding *layers*.
- Install and load the `ggplot2` package:

```
install.packages("ggplot2")  
library(ggplot2)
```

Basic Histogram:

```
ggplot(penguins, aes(body_mass_g)) +  
  geom_histogram()
```



Customizing Histograms


We can customize this with options and by adding *layers*:

- Choosing the number of bins.
- Changing the color of the bins.
- Specifying the axis labels.
- Changing the plot theme (`theme_minimal()` for removing the background colors).

```
ggplot(penguins, aes(body_mass_g)) +  
  geom_histogram(bins = 15, fill = "navy") +  
  xlab("Penguin weight (grams)") +  
  ylab("Count") +  
  theme_minimal()
```


(Output on the next slide)

Customizing Histograms



Basic Bar Plot

```
ggplot(penguins, aes(species)) +  
  geom_bar()
```



Cross-Tabulation

- We can also create a bar plot with a cross-tabulation.
- When we put 2 variables x and y in the `table()` function with `table(x, y)`, it shows us how often each combination of the values in x and y appear together in the data:


```
table(penguins$species, penguins$island)
```

	Biscoe	Dream	Torgersen
Adelie	44	56	52
Chinstrap	0	68	0
Gentoo	124	0	0

- Adelie penguins appear on all 3 islands.
- Chinstrap penguins appear only on Dream island.
- Gentoo penguins appear only on Biscoe island.

Basic Bar Plot with Cross-Tabulation

```
ggplot(penguins, aes(species, fill = island)) +  
  geom_bar()
```




Customizing Bar Plots

- We can change the name of the legend and the “fill” colors using the `scale_fill_discrete()` option.
- We can also specify colors with their *hexidecimal format* instead of color names.
 - ▶ We can find the hexadecimal format of a color using any color picker tool.

```
ggplot(penguins, aes(species, fill = island)) +  
  geom_bar(color = "black") +  
  xlab("Penguin species") +  
  ylab("Count") +  
  scale_fill_discrete(name = "Island",  
                      type = c("#0B0405", "#357BA2", "#DEF5E5")) +  
  theme_minimal()
```


(Output on the next slide)

Customizing Bar Plots




Basic Scatter Plots

```
ggplot(penguins, aes(bill_length_mm, flipper_length_mm)) +  
  geom_point()
```



Different Colors for Different Categories

```
ggplot(penguins, aes(bill_length_mm, flipper_length_mm, color = species))  
  geom_point()
```




Customizing Scatter Plots

```
ggplot(penguins, aes(bill_length_mm, flipper_length_mm, color = species))  
  geom_point() +  
  scale_color_discrete(name = "Species") +  
  xlab("Bill length (in mm)") +  
  ylab("Flipper length (in mm)") +  
  theme_minimal()
```

(Output on the next slide)

Customizing Scatter Plots



Making your own R Functions

- It's very easy to create your own R functions.
- Consider the quadratic function:

$$f(x) = -8 - 2x + x^2$$

- We can create an R function, which we call `f()`, to calculate the output of this function as follows:

```
f <- function(x) {  
  y <- -8 - 2 * x + x^2  
  return(y)  
}
```

Making your own R Functions

- We can then use this custom function like we would any other R function.
- The function evaluated at $x = 2$ should equal:

$$f(2) = -8 - 2 \times (2) + (2)^2 = -8 - 4 + 4 = -8$$

- We can check that our custom function gets the same answer:

```
f(2)
```

```
[1] -8
```

- We can also pass vectors into our custom function:

```
f(c(2, 3, 4))
```

```
[1] -8 -5  0
```


Plotting Functions with ggplot.

- We can also plot custom functions like this with ggplot.
- We first choose a range of values of x for which we want to plot the function.
- We then create a sequence of values of x in this range.
- We then evaluate the function at each of these x values to get y .
- We then put these x and y values in a `data.frame` and plot it with ggplot.

```
library(ggplot2)
x <- seq(from = -4, to = 6, length.out = 200)
y <- f(x)
df <- data.frame(x, y)
ggplot(df, aes(x, y)) + geom_line()
```

- Sometimes we don't know what range of x to choose. In this case it's good to pick some values, make the plot, and then adjust the values and make the plot again.

Plotting Functions with ggplot.



Annuity Mortgages

- The fixed amount a to be paid each period in an annuity mortgage with a loan of $\text{€}K$, a term of T periods, and a fixed interest rate of $i\%$ per period is:

$$a = K \times r^T \times \frac{1 - r}{1 - r^T}$$

where:

$$r = 1 + \frac{i}{100}$$

- For example, with a loan amount of $\text{€}5000$, a 5-year term and a 9% interest, we could use the formula to obtain:

$$a = 5000 \times (1.09)^5 \times \frac{1 - 1.09}{1 - 1.09^5} = 1285.462$$

- You would need to make 5 payments of $\text{€}1285.46$ to pay off the loan.

Annuity Mortgages in R

- A mortgage broker might need to calculate this quickly during meetings with clients.
- An R function can calculate this very quickly:

```
annuityMortgage <- function(loan, term, rate) {  
  r <- 1 + rate / 100  
  a <- loan * r^term * (1 - r) / (1 - r^term)  
  return(a)  
}
```

- 3 arguments/inputs: loan, term and rate.
- We create an R function with many inputs by separating the inputs with commas in the function() function.

```
annuityMortgage(loan = 5000, term = 5, rate = 9)
```

```
[1] 1285.462
```

Zero Points of a Quadratic Function

- The points where the function $f(x) = ax^2 + bx + c$ equals zero is given by:

$$\frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$

- The following R function can calculate these roots for any input a , b and c :

```
quadRoots <- function(a, b, c) {  
  x1 <- (-b - sqrt(b^2 - 4*a*c)) / (2*a)  
  x2 <- (-b + sqrt(b^2 - 4*a*c)) / (2*a)  
  return(c(x1, x2))  
}  
quadRoots(a = 1, b = 5, c = 4)
```

```
[1] -4 -1
```

Univariate Unconstrained Optimization

- When we plotted the function $f(x) = -8 - 2x + x^2$, we saw that it achieved a minimum at $x = 1$.
- We could solve for the minimum analytically by setting the first derivative of the function to zero:

$$\frac{df(x)}{dx} = -2 + 2x = 0 \quad \Rightarrow \quad x = 1$$

- We can also use R to find the minimum of the function using the `optimize()` function.

The optimize() Function

- We need to specify an interval (lower bound and upper bound) to search for the extreme point.
- We also need to specify if we want a maximum or a minimum using the maximum option.

```
optimize(f, interval = c(-100, 100), maximum = FALSE)
```

```
$minimum
```

```
[1] 1
```

```
$objective
```

```
[1] -9
```

- The optimize() function finds a minimum at 1, and the function takes a value of -9 at the minimum, i.e. $f(1) = -9$
- If we instead wanted to find the maximum of a function, we would specify maximum = TRUE.

The optimize() Function

- The output of the optimize() function is a list.
- If we assign the output to an object called f_min, we can extract the minimum with f_min\$minimum.
 - ▶ The \$ extraction operator works for lists just like with dataframes.
- We can similarly get the value of the function at the minimum with f_min\$objective.

```
f_min <- optimize(f, interval = c(-100, 100), maximum = FALSE)
f_min$minimum
```

```
[1] 1
```

```
f_min$objective
```

```
[1] -9
```

Basic Order Quantity Model

- You sell $d = 8000$ bikes per year. You buy them at price $p = 1000$ each.
- The holding costs are $h = 10$ per bike and ordering costs are $c = 100$ per order.
- You need to determine the number of bikes per order to minimize costs:

$$TC(q) = \underbrace{\frac{cd}{q}}_{\text{Total ordering costs}} + \underbrace{pd}_{\text{Total cost of bikes}} + \underbrace{\frac{hq}{2}}_{\text{Total holding costs}}$$

- Optimal solution is when $TC'(q) = 0$, which is $q = \sqrt{\frac{2cd}{h}}$.
- With $c = 100$, $d = 8000$ and $h = 10$, we get
 $q = \sqrt{2 \times 100 \times 8000 / 10} = 400$.

Basic Order Quantity Model in R

- We can use R to find the optimal solution without calculus. We make a function:

```
tc <- function(q) {  
  total_costs <- c*d/q + p*d + h*q/2  
  return(total_costs)  
}
```

- We set values for d, p, c and h and use optimize():

```
d <- 8000 # total demand  
p <- 1000 # price per bike  
c <- 100  # ordering costs  
h <- 10   # holding costs  
optimize(tc, c(0, 8000))
```

```
$minimum  
[1] 399.9998
```

Conditional Statements (“If-else”)

- Very often we to perform different actions depending on whether something is true or not.
- For this we use *if-else statements*, which also called *conditional statements*.
- The absolute value function is a simple example of this:

$$|x| = \begin{cases} -x & x < 0 \\ x & \text{otherwise} \end{cases}$$

- We ask, “is $x < 0$?”. If yes, then return $-x$. If not, then return x .
- Of course, we can always use the `abs( )` function in R to calculate the absolute value. But let’s create our own function doing exactly this.

Custom absolute value function

```
my_abs <- function(x) {  
  if (x < 0) {  
    return(-x)  
  } else {  
    return(x)  
  }  
}  
my_abs(-2)
```

```
[1] 2
```

```
my_abs(3)
```

```
[1] 3
```

If-Else Statements for Vectors

- The previous function we wrote only works with scalars (vectors of length 1).
- If we would try to do `my_abs(c(-2, 3))`, we would get an error.
- To do if-else statements with vectors, we can use the `ifelse()` function.
- The `ifelse()` function takes 3 arguments:
 - 1 A logical vector (such as testing the condition $x < 0$).
 - 2 What to return if TRUE (i.e. if $x < 0$).
 - 3 What to return if FALSE (i.e. if $x \geq 0$).

```
x <- -5:5  
x
```

```
[1] -5 -4 -3 -2 -1  0  1  2  3  4  5
```

```
ifelse(x < 0, -x, x)
```

```
[1] 5 4 3 2 1 0 1 2 3 4 5
```

“If-Else If-Else” Statements

- Sometimes there can be more than 2 cases to test.
- For example, consider the following function which gives the “sign” in front of a number:

$$\text{sgn}(x) = \begin{cases} -1 & \text{if } x < 0 \\ 0 & \text{if } x = 0 \\ +1 & \text{otherwise} \end{cases}$$

- For example, $\text{sgn}(-2) = -1$, $\text{sgn}(0) = 0$, and $\text{sgn}(3) = 1$.
- We can code a function in R to do this by nesting if-else statements.

“If-Else If-Else” Statements

```
sgn <- function(x) {  
  if (x < 0) {  
    return(-1)  
  } else if (x == 0) {  
    return(0)  
  } else {  
    return(+1)  
  }  
}  
sgn(-2)
```

```
[1] -1
```

```
sgn(3)
```

```
[1] 1
```

“If-Else If-Else” Statements with Vectors

We can also nest the `ifelse()` function inside of itself to get the sign of a vector of numbers:

```
x <- -3:3  
x
```

```
[1] -3 -2 -1  0  1  2  3
```

```
ifelse(x < 0, -1, ifelse(x == 0, 0, 1))
```

```
[1] -1 -1 -1  0  1  1  1
```

Merging

- Merging (or joining) is the R equivalent of the VLOOKUP function in Excel.
- When two datasets have a common ID variable linking them together, we can merge them.
- For example:
 - ▶ One dataset with total sales on each day, and another dataset with the temperature on each day. We can link the datasets using the date.
 - ▶ One dataset with the total sales in each municipality over a year, and another with the demographic characteristics of each municipality. We can link the datasets using the municipality name.

Merging Example

- We will show a merging example using 2 datasets:
 - ① Average daily petrol prices from 2014-2022.
 - ② Brent crude oil spot prices from 1987-2022 (excludes weekends and holidays).
- We will merge the datasets by date.

Data Cleaning: Average Daily Petrol Price Data

```
df1 <- read.csv("avg_daily_petrol_prices.csv")  
head(df1$date)
```

```
[1] "2014-06-08" "2014-06-09" "2014-06-10" "2014-06-11" "2014-06-12"  
[6] "2014-06-13"
```

```
df1$date <- as.Date(df1$date, format = "%Y-%m-%d")  
summary(df1)
```

date	e5	e10	diesel
Min. :2014-06-08	Min. :1.159	Min. :1.130	Min. :0.9558
1st Qu.:2016-07-03	1st Qu.:1.340	1st Qu.:1.318	1st Qu.:1.1322
Median :2018-07-29	Median :1.402	Median :1.379	Median :1.2353
Mean :2018-07-29	Mean :1.456	Mean :1.423	Mean :1.2811
3rd Qu.:2020-08-23	3rd Qu.:1.522	3rd Qu.:1.479	3rd Qu.:1.3217
Max. :2022-09-18	Max. :2.261	Max. :2.203	Max. :2.3343

Data Cleaning: Brent Crude Oil Spot Price Data

```
df2 <- read.csv("Europe_Brent_Spot_Price_FOB.csv", skip = 4)
head(df2$Day)
```

```
[1] "09/19/2022" "09/16/2022" "09/15/2022" "09/14/2022" "09/13/2022"
[6] "09/12/2022"
```

```
df2$Day <- as.Date(df2$Day, format = "%m/%d/%Y")
names(df2) <- c("date", "crude_oil")
summary(df2)
```

date	crude_oil
Min. :1987-05-20	Min. : 9.10
1st Qu.:1996-03-06	1st Qu.: 19.03
Median :2005-01-04	Median : 38.08
Mean :2005-01-12	Mean : 48.22
3rd Qu.:2013-11-24	3rd Qu.: 69.67
Max. :2022-09-19	Max. :143.95

The merge() Command:

```
df <- merge(df1, df2, by = "date")
```

```
summary(df)
```

date	e5	e10	diesel
Min. :2014-06-09	Min. :1.159	Min. :1.130	Min. :0.9558
1st Qu.:2016-07-04	1st Qu.:1.339	1st Qu.:1.318	1st Qu.:1.1302
Median :2018-07-26	Median :1.402	Median :1.379	Median :1.2345
Mean :2018-07-27	Mean :1.455	Mean :1.422	Mean :1.2797
3rd Qu.:2020-08-19	3rd Qu.:1.521	3rd Qu.:1.478	3rd Qu.:1.3216
Max. :2022-09-16	Max. :2.261	Max. :2.203	Max. :2.3343

crude_oil
Min. : 9.12
1st Qu.: 48.54
Median : 61.18
Mean : 63.47
3rd Qu.: 72.97
Max. :133.18

Merging: Dropped Observations

- The merged dataset only includes observations where there is a match.
- Observations where there is no corresponding match are dropped:

```
nrow(df1)
```

```
[1] 3025
```

```
nrow(df2)
```

```
[1] 8970
```

```
nrow(df)
```

```
[1] 2107
```

- To avoid dropping rows, we can use the `all.x = TRUE` and/or `all.y = TRUE` options.

Merging with `all.x = TRUE`

- `all.x = TRUE` : Keeps all observations in the 1st dataset, but only merges data from the 2nd dataset when there is a match. When there is no match, variables in the 2nd dataset get assigned NA values.
- `all.y = TRUE` : Keeps all observations in the 2nd dataset, but only merges data from the 1st dataset when there is a match. When there is no match, variables in the 1st dataset get assigned NA values.
- `all = TRUE` : This keeps all observations from both datasets, and variables get assigned NA values when there is no match. This is equivalent to setting both `all.x = TRUE` and `all.y = TRUE`.

Merging with `all.x = TRUE`

```
df <- merge(df1, df2, by = "date", all.x = TRUE)
```

```
summary(df)
```

date		e5	e10	diesel
Min.	:2014-06-08	Min. :1.159	Min. :1.130	Min. :0.9558
1st Qu.:	:2016-07-03	1st Qu.:1.340	1st Qu.:1.318	1st Qu.:1.1322
Median	:2018-07-29	Median :1.402	Median :1.379	Median :1.2353
Mean	:2018-07-29	Mean :1.456	Mean :1.423	Mean :1.2811
3rd Qu.:	:2020-08-23	3rd Qu.:1.522	3rd Qu.:1.479	3rd Qu.:1.3217
Max.	:2022-09-18	Max. :2.261	Max. :2.203	Max. :2.3343

crude_oil

Min.	: 9.12
1st Qu.:	48.54
Median	: 61.18
Mean	: 63.47
3rd Qu.:	72.97
Max.	:133.18
NAs	:918

Further Remarks on Merging

- If merging on multiple variables, use `by = c("var1", "var2")`.
- If merging variables differ in `df1` and `df2`, use `by.x` and `by.y` instead of `by`. For example:

```
df <- merge(df1, df2, by.x = c("market_area", "date"),  
            by.y = c("market", "date"))
```

- To avoid sorting the data, use option `sort = FALSE`.

Reshaping

Suppose you have a dataset structured like this (long format):

	id	variable	value
1	1	x	3
2	1	y	5
3	2	x	4
4	2	y	8
5	3	x	3
6	3	y	1

And you wanted to *reshape* it to look like this (wide format):

	id	x	y
1	1	3	5
2	2	4	8
3	3	3	1

Reshaping

- Base R has a function that can do this called `reshape()`, but it's not very easy to use.
- The package `reshape2` contains functions to make it easy to go from long to wide format and vice-versa:
 - ▶ `dcast(df, id ~ variable, value.var = "value")`: long to wide format.
 - ▶ `melt(df, idvars = "id")`: wide to long format.

Reshaping: Long to Wide

```
long <- data.frame(
  id      = rep(1:3, each = 2),
  variable = rep(c("x", "y"), times = 3),
  value   = c(3, 5, 4, 8, 3, 1)
)
library(reshape2)
wide <- dcast(long, id ~ variable, value.var = "value")
wide
```

	id	x	y
1	1	3	5
2	2	4	8
3	3	3	1

Reshaping: Wide to Long

```
melt(wide, id.vars = "id")
```

	id	variable	value
1	1	x	3
2	2	x	4
3	3	x	3
4	1	y	5
5	2	y	8
6	3	y	1

Overlay Line Plots with ggplot

To overlay line plots with ggplot we need our data in long format.

```
df <- read.csv("avg_daily_petrol_prices.csv")
df$date <- as.Date(df$date, format = "%Y-%m-%d")
head(df, n = 2)
```

	date	e5	e10	diesel
1	2014-06-08	1.551987	1.477774	1.353583
2	2014-06-09	1.576623	1.483362	1.385182

```
df2 <- melt(df, id.vars = "date")
head(df2, n = 2)
```

	date	variable	value
1	2014-06-08	e5	1.551987
2	2014-06-09	e5	1.576623

Overlay Line Plots with ggplot

```
ggplot(df2, aes(date, value, color = variable)) + geom_line()
```




Customizing the Plot

```
levels(df2$variable) <- c("E5", "E10", "Diesel")
library(ggplot2)
ggplot(df2, aes(date, value, color = variable)) +
  geom_line() +
  xlab("") +
  ylab("Average Daily\nPetrol Price\n(in Euro)") +
  scale_color_discrete(name = "Petrol Type:") +
  theme_minimal() +
  theme(legend.direction = "horizontal",
        legend.position = "bottom")
```

(Output on next slide)

Customizing the Plot



Aggregating by Group

- If we want to get the sum or average by group we can use the `aggregate()` function.
- The `aggregate()` function is a bit like the R version of pivot tables in Excel.
- If we want to get the average of `x` by group `g` in the dataframe `df`, we use:

```
aggregate(x ~ g, FUN = mean, data = df)
```

- We'll show this using the petrol price data as an example.

Average Price of E5 Petrol by Year

```
df <- read.csv("avg_daily_petrol_prices.csv")
df$date <- as.Date(df$date, format = "%Y-%m-%d")
library(lubridate)
df$year <- year(df$date)
aggregate(e5 ~ year, FUN = mean, data = df)
```

	year	e5
1	2014	1.519195
2	2015	1.393170
3	2016	1.302767
4	2017	1.368414
5	2018	1.454098
6	2019	1.430592
7	2020	1.288080
8	2021	1.579992
9	2022	1.933512

Maximum Price of E5 by Year

```
aggregate(e5 ~ year, FUN = max, data = df)
```

	year	e5
1	2014	1.605252
2	2015	1.519229
3	2016	1.395620
4	2017	1.414228
5	2018	1.573760
6	2019	1.555623
7	2020	1.461060
8	2021	1.761985
9	2022	2.260581

Average Price of E5 and E10 by Year

```
aggregate(cbind(e5, e10) ~ year, FUN = mean, data = df)
```

	year	e5	e10
1	2014	1.519195	1.461632
2	2015	1.393170	1.373425
3	2016	1.302767	1.282375
4	2017	1.368414	1.345430
5	2018	1.454098	1.430955
6	2019	1.430592	1.408177
7	2020	1.288080	1.253603
8	2021	1.579992	1.522833
9	2022	1.933512	1.875897

Average Price of All Variables by Year

```
aggregate(. ~ year, FUN = mean, data = df)
```

	year	date	e5	e10	diesel
1	2014	16332.0	1.519195	1.461632	1.334612
2	2015	16618.0	1.393170	1.373425	1.173013
3	2016	16983.5	1.302767	1.282375	1.081282
4	2017	17349.0	1.368414	1.345430	1.161306
5	2018	17714.0	1.454098	1.430955	1.287264
6	2019	18079.0	1.430592	1.408177	1.265341
7	2020	18444.5	1.288080	1.253603	1.111068
8	2021	18810.0	1.579992	1.522833	1.387114
9	2022	19123.0	1.933512	1.875897	1.941386